

References

1. DBpedia. <http://wiki.dbpedia.org/>
2. Freebase. <https://www.freebase.com/>
3. UniProt. <ftp://ftp.uniprot.org/>
4. Abadi, D.J., Marcus, A., Madden, S.R., Hollenbach, K.: Scalable semantic web data management using vertical partitioning. In: VLDB. pp. 411–422 (2007)
5. Abadi, D.J., Marcus, A., Madden, S.R., Hollenbach, K.: SW-Store: A vertically partitioned DBMS for semantic web data management. PVLDB **18**(2), 385–406 (2009)
6. Abdel Kader, R., Boncz, P., Manegold, S., van Keulen, M.: ROX: Run-time optimization of XQueries. In: SIGMOD. pp. 615–626 (2009)
7. Bornea, M., Dolby, J., Kementsietsidis, A., Srinivas, K., Dantressangle, P., Udrea, O., Bhattacharjee, B.: Building an efficient RDF store over a relational database. In: SIGMOD. pp. 121–132 (2013)
8. Broekstra, J., Kampman, A., Harmelen, F.v.: Sesame: A generic architecture for storing and querying RDF and RDF schema. In: ISWC. pp. 54–68 (2002)
9. Bruno, N., Chaudhuri, S.: Exploiting statistics on query expressions for optimization. In: SIGMOD. pp. 263–274 (2002)
10. Chaudhuri, S.: An overview of query optimization in relational systems. In: PODS. pp. 34–43 (1998)
11. Chong, E.I., Das, S., Eadon, G., Srinivasan, J.: An efficient SQL-based RDF querying scheme. In: VLDB. pp. 1216–1227 (2005)
12. Cole, R.L., Graefe, G.: Optimization of dynamic query evaluation plans. In: SIGMOD. pp. 150–160 (1994)
13. Eric, P., Alexandre, B.: A mapping of SPARQL onto conventional SQL. <https://www.w3.org/2008/07/MappingRules/StemMapping> (2008)
14. Erling, O., Mikhailov, I.: Virtuoso: RDF support in a native RDBMS. In: Semantic Web Information Management, pp. 501–519. Springer (2009)
15. Graefe, G., Ward, K.: Dynamic query evaluation plans. In: SIGMOD. pp. 358–366 (1989)
16. Gubichev, A., Neumann, T.: Exploiting the query structure for efficient join ordering in SPARQL queries. In: EDBT. pp. 439–450 (2014)
17. Harris, S., Seaborne, A.: SPARQL 1.1 query language. W3C working draft. <http://www.w3.org/TR/sparql11-query> (2013)
18. Harth, A., Hose, K., Schenkel, R.: Database techniques for linked data management. In: SIGMOD. pp. 597–600 (2012)
19. Ioannidis, Y.E.: The history of histograms (abridged). In: VLDB. pp. 19–30 (2003)
20. Ioannidis, Y.E., Christodoulakis, S.: On the propagation of errors in the size of join results. In: SIGMOD. pp. 268–277 (1991)
21. Ioannidis, Y.E., Christodoulakis, S.: Optimal histograms for limiting worst-case error propagation in the size of join results. ACM Trans. Database Syst. **18**(4), 709–748 (Dec 1993)
22. Kotoulas, S., Urbani, J., Boncz, P., Mika, P.: Robust runtime optimization and skew-resistant execution of analytical sparql queries on pig. In: ISWC. pp. 247–262 (2012)
23. Leis, V., Gubichev, A., Mirchev, A., Boncz, P., Kemper, A., Neumann, T.: How good are query optimizers, really? PVLDB **9**(3) (Nov 2015)
24. Lipton, R.J., Naughton, J.F., Schneider, D.A.: Practical selectivity estimation through adaptive sampling. In: SIGMOD. pp. 1–11 (1990)

25. Lohman, G.: Is query optimization a solved problem? <http://wp.sigmod.org/?p=1075> (2014)
26. Neumann, T., Galindo-Legaria, C.A.: Taking the edge off cardinality estimation errors using incremental execution. In: BTW. pp. 73–92 (2013)
27. Neumann, T., Moerkotte, G.: Characteristic sets: Accurate cardinality estimation for RDF queries with multiple joins. In: ICDE. pp. 984–994 (2011)
28. Neumann, T., Weikum, G.: The RDF-3X engine for scalable management of RDF data. In: PVLDB. vol. 19, pp. 91–113 (2010)
29. Pérez, J., Arenas, M., Gutierrez, C.: Semantics and complexity of SPARQL. *ACM Trans. Database Syst.* **34**(3), 16:1–16:45 (2009)
30. Pham, M., Boncz, P.A., Erling, O.: S3G2: A scalable structure-correlated social graph generator. In: TPCTC. pp. 156–172 (2012)
31. Picalausa, F., Vansummeren, S.: What are real SPARQL queries like? In: SWIM. pp. 7:1–7:6 (2011)
32. Schmidt, M., Meier, M., Lausen, G.: Foundations of SPARQL query optimization. In: ICDT. pp. 4–33 (2010)
33. Stocker, M., Seaborne, A., Bernstein, A., Kiefer, C., Reynolds, D.: SPARQL basic graph pattern optimization using selectivity estimation. In: WWW. pp. 595–604 (2008)
34. Tsialiamanis, P., Sidiropoulos, L., Fundulaki, I., Christophides, V., Boncz, P.: Heuristics-based query optimisation for SPARQL. In: EDBT. pp. 324–335 (2012)
35. Weiss, C., Karras, P., Bernstein, A.: Hexastore: Sextuple indexing for semantic web data management. *PVLDB*. **1**(1) (Aug 2008)
36. Wilkinson, K.: Jena Property Table Implementation. In: SSWS. pp. 35–46 (2006)
37. Yuan, P., Liu, P., Wu, B., Jin, H., Zhang, W., Liu, L.: TripleBit: A fast and compact system for large scale RDF data. *PVLDB* **6**(7), 517–528 (2013)